

Movie Reviews Sentiment Analysis with Scikit-Learn

PyLing meeting, Feb 8 2017. Updated Feb 21, 2019

Following this tutorial on scikit-learn.org: http://scikit-learn.org/dev/tutorial/text_analytics/working_with_text_data.html (http://scikit-learn.org/dev/tutorial/text_analytics/working_with_text_data.html)

Adapted to

- work with off-line movie review corpus, which was also covered/used in NLTK book (<http://www.nltk.org/book/ch06.html#document-classification>), downloadable here (http://www.nltk.org/nltk_data/).
- use the NLTK's tokenizer (so symbols and stopwords are not thrown out)

Also, check out documentation on dataset loading: <https://scikit-learn.org/stable/datasets.html> (<https://scikit-learn.org/stable/datasets.html>)

Load movie_reviews corpus data through sklearn

```
In [1]: import sklearn
        from sklearn.datasets import load_files
```

```
In [2]: moviedir = r'D:\Lab\nltk_data\corpora\movie_reviews'

        # loading all files.
        movie = load_files(moviedir, shuffle=True)
```

```
In [3]: len(movie.data)

Out[3]: 2000
```

```
In [4]: # target names ("classes") are automatically generated from subfolder names
        movie.target_names
```

```
Out[4]: ['neg', 'pos']
```

```
In [5]: # First file seems to be about a Schwarzenegger movie.  
movie.data[0][:500]
```

```
Out[5]: b"arnold schwarzenegger has been an icon for action enthusiasts ,  
since the late 80's , but lately his films have been very sloppy a  
nd the one-liners are getting worse . \nit's hard seeing arnold as  
mr . freeze in batman and robin , especially when he says tons of  
ice jokes , but hey he got 15 million , what's it matter to him ?  
\nonce again arnold has signed to do another expensive blockbuster  
, that can't compare with the likes of the terminator series , tru  
e lies and even eraser . \nin this so cal"
```

```
In [6]: # first file is in "neg" folder  
movie.filenames[0]
```

```
Out[6]: 'D:\\Lab\\nltk_data\\corpora\\movie_reviews\\neg\\cv405_21868.txt'
```

```
In [7]: # first file is a negative review and is mapped to 0 index 'neg' in  
target_names  
movie.target[0]
```

```
Out[7]: 0
```

A detour: try out CountVectorizer & TF-IDF

```
In [8]: # import CountVectorizer, nltk  
from sklearn.feature_extraction.text import CountVectorizer  
import nltk
```

```
In [9]: # Turn off pretty printing of jupyter notebook... it generates long  
lines  
%pprint
```

```
Pretty printing has been turned OFF
```

```
In [10]: # Three tiny "documents"  
docs = ['A rose is a rose is a rose is a rose.',  
       'Oh, what a fine day it is.',  
       "A day ain't over till it's truly over."]
```

```
In [11]: # Initialize a CountVectorizer to use NLTK's tokenizer instead of its  
# default one (which ignores punctuation and stopwords).  
# Minimum document frequency set to 1.  
fooVzer = CountVectorizer(min_df=1, tokenizer=nltk.word_tokenize)
```

```
In [12]: # .fit_transform does two things:  
# (1) fit: adapts fooVzer to the supplied text data (rounds up top  
words into vector space)  
# (2) transform: creates and returns a count-vectorized output of d  
ocs  
docs_counts = fooVzer.fit_transform(docs)  
  
# fooVzer now contains vocab dictionary which maps unique words to  
indexes  
fooVzer.vocabulary_
```

```
Out[12]: {'a': 3, 'rose': 12, 'is': 7, '.': 2, 'oh': 10, ',': 1, 'what': 1  
5, 'fine': 6, 'day': 5, 'it': 8, 'ai': 4, "n't": 9, 'over': 11, 't  
ill': 13, "'s": 0, 'truly': 14}
```

```
In [13]: # docs_counts has a dimension of 3 (document count) by 16 (# of uni  
que words)  
docs_counts.shape
```

```
Out[13]: (3, 16)
```

```
In [14]: # this vector is small enough to view in a full, non-sparse form!  
docs_counts.toarray()
```

```
Out[14]: array([[0, 0, 1, 4, 0, 0, 0, 3, 0, 0, 0, 0, 4, 0, 0, 0],  
[0, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 1],  
[1, 0, 1, 1, 1, 1, 0, 0, 1, 1, 0, 2, 0, 1, 1, 0]], dtype=in  
t64)
```

```
In [15]: # Convert raw frequency counts into TF-IDF (Term Frequency -- Inver  
se Document Frequency) values  
from sklearn.feature_extraction.text import TfidfTransformer  
fooTfmer = TfidfTransformer()  
  
# Again, fit and transform  
docs_tfidf = fooTfmer.fit_transform(docs_counts)
```

```
In [16]: # TF-IDF values
# raw counts have been normalized against document length,
# terms that are found across many docs are weighted down ('a' vs.
'rose')
docs_tfidf.toarray()

Out[16]: array([[0.          , 0.          , 0.11337964, 0.45351858, 0.
,
0.          , 0.          , 0.4379908 , 0.          , 0.
,
0.          , 0.          , 0.7678737 , 0.          , 0.
,
0.          ],
[0.          , 0.39427404, 0.2328646 , 0.2328646 , 0.
,
0.29985557, 0.39427404, 0.29985557, 0.29985557, 0.
,
0.39427404, 0.          , 0.          , 0.          , 0.
,
0.39427404],
[0.30352608, 0.          , 0.17926739, 0.17926739, 0.3035260
8,
0.23083941, 0.          , 0.          , 0.23083941, 0.3035260
8,
0.          , 0.60705216, 0.          , 0.30352608, 0.3035260
8,
0.          ]])
```

So that completes the feature extraction process for our toy document set

- What if we get new ones?

```
In [17]: # A list of new documents
newdocs = ["I have a rose and a lily.", "What a beautiful day."]

# This time, no fitting needed: transform the new docs into count-vectorized form
# Unseen words ('lily', 'beautiful', 'have', etc.) are ignored
newdocs_counts = fooVzer.transform(newdocs)
newdocs_counts.toarray()

Out[17]: array([[0, 0, 1, 2, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0],
[0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1]], dtype=in
t64)
```

```
In [18]: # Again, transform using tfidf  
newdocs_tfidf = fooTfmer.transform(newdocs_counts)  
newdocs_tfidf.toarray()
```

```
Out[18]: array([[0.          , 0.          , 0.35653519, 0.71307037, 0.  
,  
     0.          , 0.          , 0.          , 0.          , 0.  
,  
     0.          , 0.          , 0.60366655, 0.          , 0.  
,  
     0.          , 0.          , 0.          , 0.          , 0.  
,  
     0.50410689, 0.          , 0.          , 0.          , 0.  
,  
     0.          , 0.          , 0.          , 0.          , 0.  
,  
     0.66283998]])
```

Back to real data: movie reviews

```
In [19]: # Split data into training and test sets  
from sklearn.model_selection import train_test_split  
docs_train, docs_test, y_train, y_test = train_test_split(movie.data,  
                                         movie.target,  
                                         test_size  
                                         = 0.20, random_state = 12)
```

```
In [20]: # initialize CountVectorizer  
movieVzer= CountVectorizer(min_df=2, tokenizer=nltk.word_tokenize,  
max_features=3000) # use top 3000 words only. 78.25% acc.  
# movieVzer = CountVectorizer(min_df=2, tokenizer=nltk.word_tokenize)  
# use all 25K words. Higher accuracy  
  
# fit and tranform using training text  
docs_train_counts = movieVzer.fit_transform(docs_train)
```

```
In [21]: # 'screen' is found in the corpus, mapped to index 2290  
movieVzer.vocabulary_.get('screen')
```

```
Out[21]: 2290
```

```
In [22]: # Likewise, Mr. Steven Seagal is present...  
movieVzer.vocabulary_.get('seagal')
```

```
Out[22]: 2297
```

```
In [23]: # huge dimensions! 1,600 documents, 3K unique terms.  
docs_train_counts.shape
```

```
Out[23]: (1600, 3000)
```

```
In [24]: # Convert raw frequency counts into TF-IDF values
movieTfmer = TfidfTransformer()
docs_train_tfidf = movieTfmer.fit_transform(docs_train_counts)
```

```
In [25]: # Same dimensions, now with tf-idf values instead of raw frequency
counts
docs_train_tfidf.shape
```

```
Out[25]: (1600, 3000)
```

The feature extraction functions and training data are ready.

- Vectorizer and transformer have been built from the training data
- Training data text was also turned into TF-IDF vector form

Next up: test data

- You have to prepare the test data using the same feature extraction scheme.

```
In [26]: # Using the fitted vectorizer and transformer, transform the test da
ta
docs_test_counts = movieVzer.transform(docs_test)
docs_test_tfidf = movieTfmer.transform(docs_test_counts)
```

Training and testing a Naive Bayes classifier

```
In [27]: # Now ready to build a classifier.
# We will use Multinomial Naive Bayes as our model
from sklearn.naive_bayes import MultinomialNB
```

```
In [28]: # Train a Multimoda Naive Bayes classifier. Again, we call it "fitt
ing"
clf = MultinomialNB()
clf.fit(docs_train_tfidf, y_train)
```

```
Out[28]: MultinomialNB(alpha=1.0, class_prior=None, fit_prior=True)
```

```
In [29]: # Predict the Test set results, find accuracy
y_pred = clf.predict(docs_test_tfidf)
sklearn.metrics.accuracy_score(y_test, y_pred)
```

```
Out[29]: 0.7825
```

```
In [30]: # Making the Confusion Matrix
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)
cm
```

```
Out[30]: array([[164,  42],
   [ 45, 149]], dtype=int64)
```

Trying the classifier on fake movie reviews

```
In [31]: # very short and fake movie reviews
reviews_new = ['This movie was excellent', 'Absolute joy ride',
               'Steven Seagal was terrible', 'Steven Seagal shone through.',
               'This was certainly a movie', 'Two thumbs up', 'I fell asleep halfway through',
               "We can't wait for the sequel!!", '!', '?', 'I cannot recommend this highly enough',
               'instant classic.', 'Steven Seagal was amazing. His performance was Oscar-worthy.']

reviews_new_counts = movieVzer.transform(reviews_new)           # turn text into count vector
reviews_new_tfidf = movieTfmer.transform(reviews_new_counts)  # turn n into tfidf vector
```

```
In [32]: # have classifier make a prediction
pred = clf.predict(reviews_new_tfidf)
```

```
In [33]: # print out results
for review, category in zip(reviews_new, pred):
    print('%r => %s' % (review, movie.target_names[category]))

'This movie was excellent' => pos
'Absolute joy ride' => pos
'Steven Seagal was terrible' => neg
'Steven Seagal shone through.' => neg
'This was certainly a movie' => neg
'Two thumbs up' => neg
'I fell asleep halfway through' => neg
"We can't wait for the sequel!!" => neg
'!' => neg
'?' => neg
'I cannot recommend this highly enough' => pos
'instant classic.' => pos
'Steven Seagal was amazing. His performance was Oscar-worthy.' =>
neg
```

```
In [34]: # Mr. Seagal simply cannot win!
```

Final notes

- In practice, you should use **TfidfVectorizer**, which is CountVectorizer and TfidfTransformer conveniently rolled into one:

```
from sklearn.feature_extraction.text import TfidfVectorizer
```

- Also: It is a popular practice to use **pipeline**, which pairs up your feature extraction routine with your choice of ML model:

```
model = make_pipeline(TfidfVectorizer(), MultinomialNB())
```